


MCTS (Checkers) AI Discussion

Disclaimer : We discuss a few ideas, there are many others... We will not go into detailed technical solutions (actual code) ...

Kalev Kask



Some pictures from publications referenced

Resources

- “A Survey of Monte Carlo Tree Search Methods”
 - http://ccg.doc.gold.ac.uk/ccg_old/papers/browne_tciaig12_1.pdf
- many others
-

Basics

- Board size
 - M = number of rows
 - N = number of columns
- StudentAI class
 - This is where your code will go ...
- `get_move()/GetMove()`
 - called by the system, when it is your turn to move ...
- Time limit
 - E.g. 8 min per game per player ...

Persistent Board

- There is a board (member) object (of type Board) in StudentAI class, for your convenience
- You need to update it yourself
 - e.g. when `get_move()` is called, opponent's move is passed in as input, you need to call `make_move()` to keep the board up to date
 - ...

Keeping track of time

- Time is of the essence

- There will be a timeout

- You have to keep track of time

- Avg # of moves per game ? 25

- 8min = 480sec -> about 20 sec per move
- MCTS iterations slower early, faster later

```
double total_time_elapsed = 0.0
```

```
...
```

```
Move get_move(...)
```

```
{
```

```
    double remaining_time = ...
```

```
    if (remaining_time < some_small_number_eg_3)
```

```
        make_random_move
```

```
    else {
```

```
        tS = time_stamp_now
```

```
        // do your normal stuff
```

```
        .....
```

```
        tE = time_stamp_now
```

```
        dt = tE - tS // time used for this get_move() call
```

```
        total_time_elapsed += dt
```

```
    }
```

```
}
```

MCTS

- Basic question :
 - How long does each iteration take
 - this determines how many iterations you can do
 - What do you get out of each iteration
- If # iterations $\rightarrow \infty$ then learned win-rate \rightarrow true win-rate
 - The more iterations you can do, the better quality play
 - Are you doing enough iterations?

MCTS tradeoff

- Make each iteration faster?
- Get more out of each iteration?
 - Only “good quality” iterations really contribute
 - Naïve MCTS will execute many “poor quality” iterations -> slow convergence
- Tradeoff : is the extra effort you put into each iterations (e.g. heuristic) paying off

MCTS how many iterations?

- How many iterations is enough?
 - Depends of what you get out of each iteration
 - For checkers
 - 100 probably not enough
 - 1000 probably enough

MCTS how to make iterations faster ?

- When doing selection/simulation you need the board along the way
 - For blind/uninformed simulation, can make a copy of “main” board in the beginning, and then just call `getAllMoves()/makeMove()`
 - If you heuristic, need to do 1-move lookahead at each state (simulation)
 - Do you do (deep)copy of the board, and then just `makeMove()` off of the (deep)copy?
 - Do you skip (deep)copy and do `makeMove()/Undo()`?

MCTS how to make iterations faster ?

- When entering `getMove()`
 - `getAllPossibleMoves()` returns just 1 single move, don't need to be MCTS
 - Just immediately return that 1 single move

MCTS how to make iterations faster ?

- When simulation is taking too long?
 - You can cut it off, but the terminal value is unknown
 - Can apply heuristic
 - But heuristic needs to be sufficiently accurate
 - May need 3 states
 - Win
 - Loss
 - TooHardToTell


Heuristic

- Heuristic needs to be
 - Accurate
 - Discriminative
- E.g. $H(s) = \#B(s) - \#W(s)$ is not discriminative
 - Vast majority of time, either no capture (H is same for all children) or 1-piece capture (H is same for all children)
 - Most of the time material value does not change
 - So this H is uninformative and useless for guiding MCTS

Heuristic : testing

- Experimentally evaluate how discriminative/accurate it is
- Implement simple MiniMax (a few moves lookahead) and run your H against Random/Poor/Average. What is your winrate?

MCTS reuse subtree



MCTS modified UCT

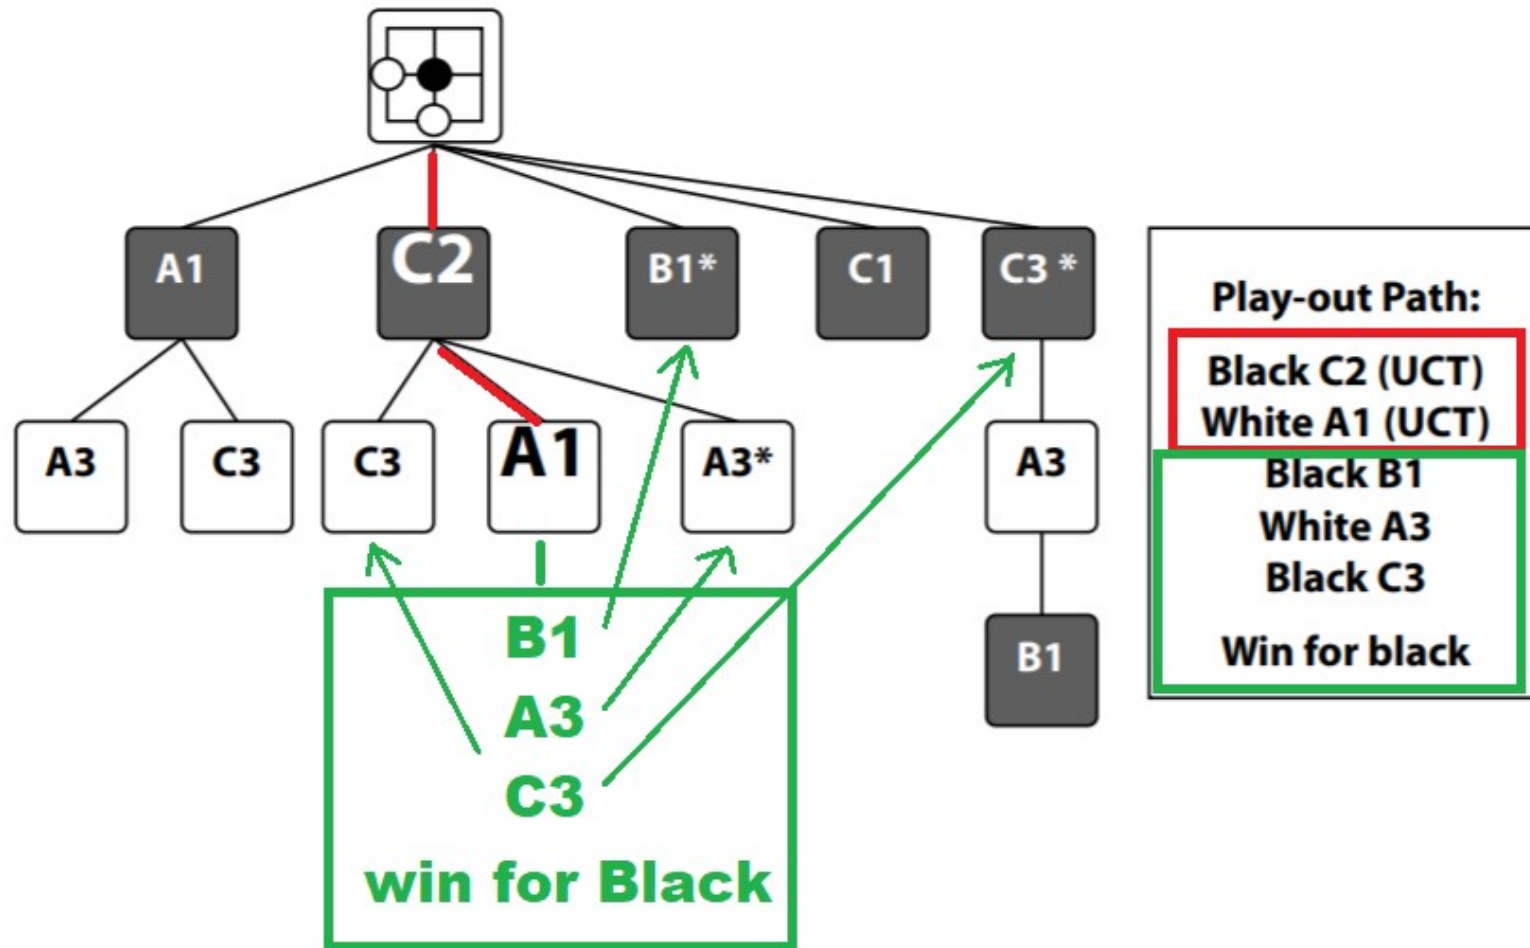
$$\text{UCT} = \frac{w_i}{s_i} + c \sqrt{\frac{\ln s_p}{s_i}} + \frac{b_i}{s_i}$$

when adding node to MCTS tree can initialize w_i and s_i

what is c ? can determine experimentally

add a term $\frac{b_i}{s_i}$


MCTS : AMAF heuristic



MCTS : AMAF heuristic

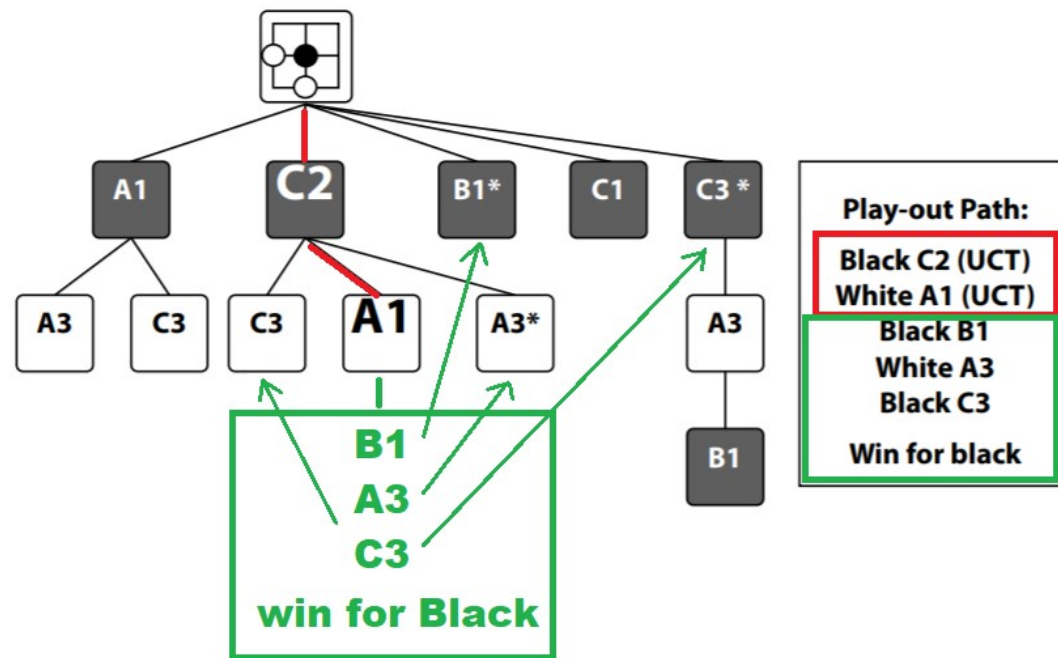
- AMAF = All Moves As First
- Value of making move **a** immediately is **average outcome of all simulations** where **a** is played **at any time**
 - Have a general value for each move, **regardless** of when it is played
- Biased estimate of true move value
- Based on independence assumption – value of move is unaffected by other moves before/after
- Advantage = quick estimates
- Disadvantage = biased

MCTS : AMAF nodes values



- For nodes on the selection path,
- Update those siblings of the selection path,
- That correspond to moves selected during simulation

- Update siblings of **Start, C2, A1** based on the path Start → C2 → A1 → B1 → A3 → C3



MCTS : α -AMAF

- Compute/keep 2 sets of counts for each node
 - Standard MCTS estimates
 - AMAF estimates
- α -AMAF estimate is $\alpha \cdot \text{AMAF} + (1-\alpha) \cdot \text{Standard}$
 - $\alpha=0$ is Standard MCTS
 - $\alpha=1$ is (pure) AMAF
 - α -AMAF is blend of the two

MCTS : RAVE

- RAVE = Rapid Action Value Estimates
- Each node in MCTS tree has its own α that starts at 1 and then decreases to 0, as more iterations go through the node
 - If you have very few sampled paths through the node, you rely on its AMAF estimate
 - If you have many sampled paths through the node, you rely on its own Standard MCTS estimate

$$UCT_{RAVE}(s) = \alpha(s) \cdot UCT_{AMAF}(s) + (1 - \alpha(s)) \cdot UCT_{Standard}(s)$$

$$\alpha(s) = \max(0, \frac{(P - s_i)}{P})$$

- P is a parameter, s_i is count for node S